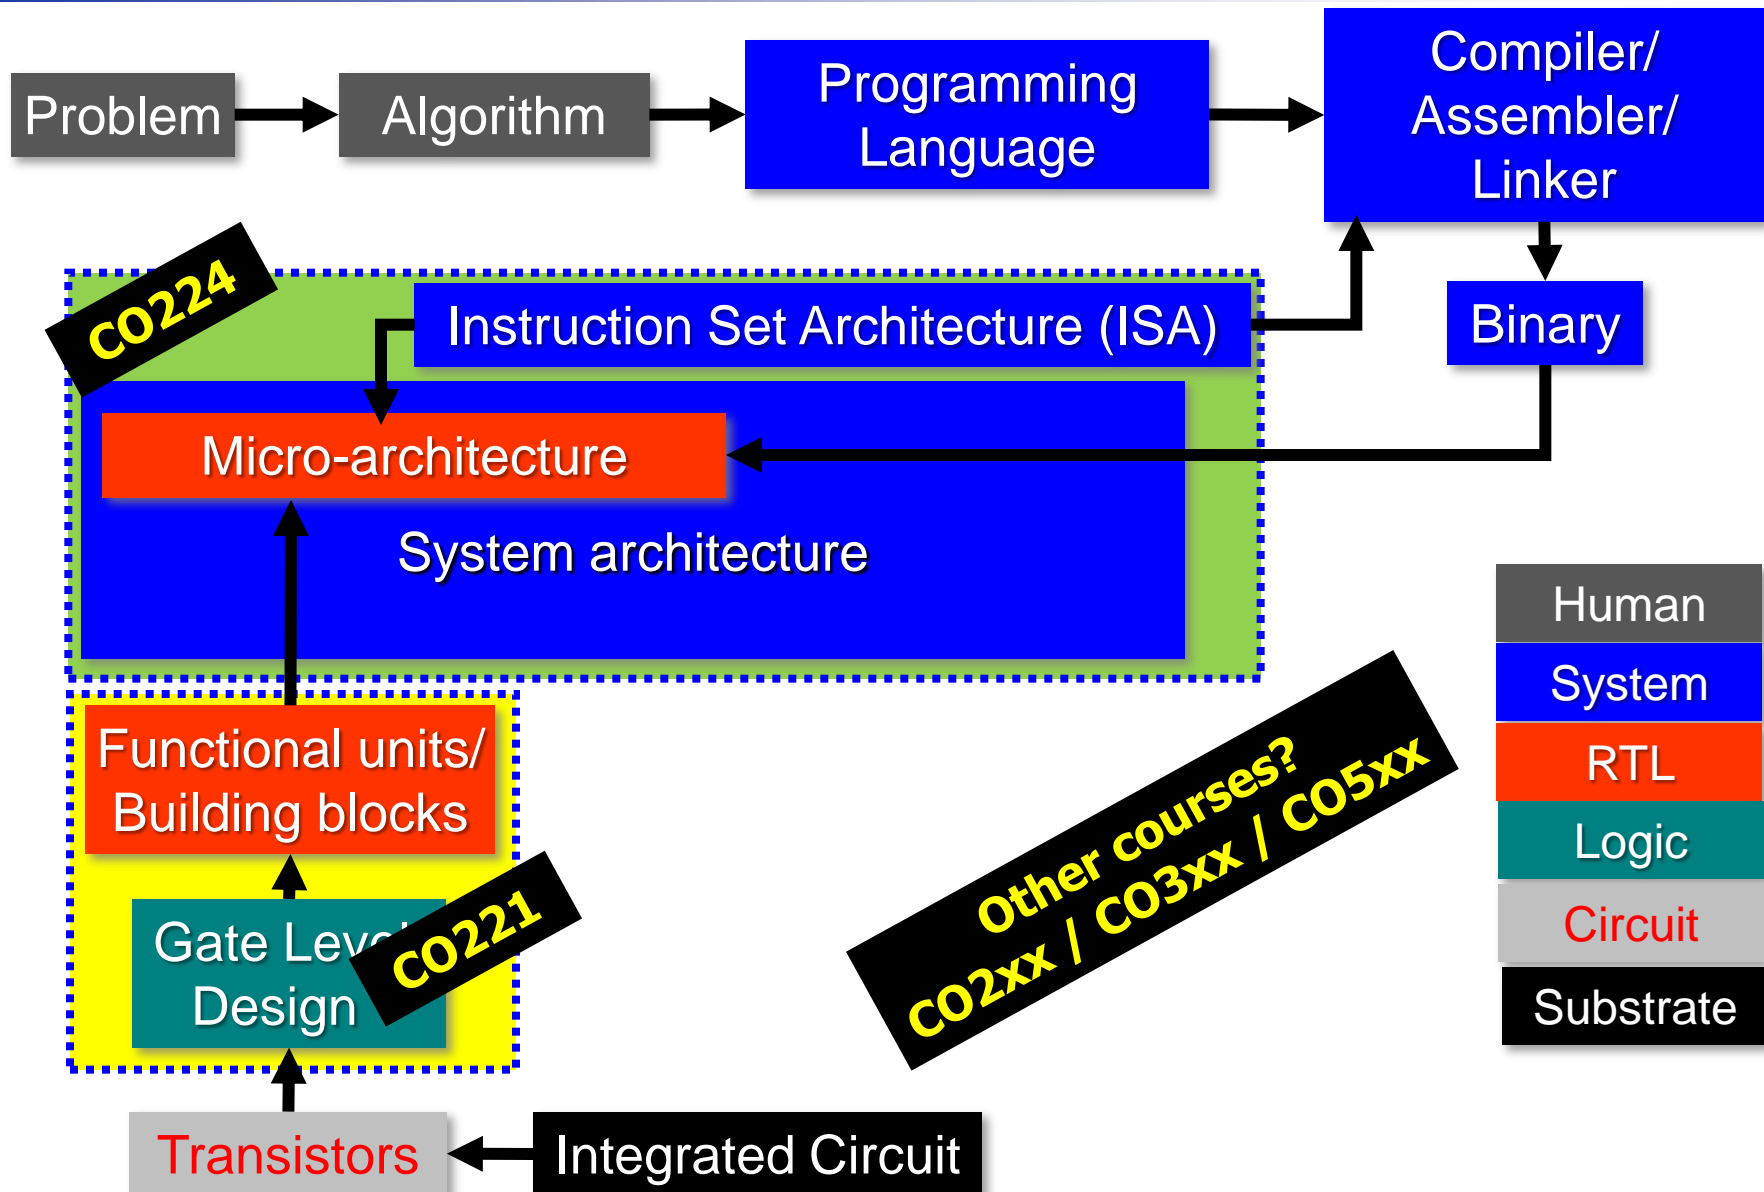


CO224: COMPUTER ARCHITECTURE

Isuru Nawinne

[Adapted from Computer Organization and Design, ARM Edition. Patterson & Hennessy, © 2011, MK]

The BIG Picture



CO224: Outline

Lectures: Computer Organization and Design

Ch01 – Computer Abstraction and Technology

Ch02 – Instructions – The Language of the Computer

Ch03 – Arithmetic for Computers

Ch04 – The Processor

Ch05 – Memory Hierarchy

Ch06 – Storage and Other IO

Ch07 – Multi-cores, Multi-processors and Clusters

Labs:

(1) Writing Assembly programs for ARM ISA

(2) Micro-architecture & systems design

What You Will Learn

- How programs are translated into the machine language
 - And how the hardware executes them
- The hardware/software interface
- What determines program performance
 - And how it can be improved
- How hardware designers improve performance
- What is parallel processing



CHAPTER 01

COMPUTER ABSTRACTIONS AND TECHNOLOGY

- **Computer Abstractions**
- Technology Trends
- The Power Wall & Multi-processors
- Below Your Program
- Under the Covers
- Performance

Computer Abstractions

- What is abstraction ?!?
 -
- Can you figure out the order of these abstractions?
 - Organize them in a meaningful way...



CHAPTER 01

COMPUTER ABSTRACTIONS AND TECHNOLOGY

- Computer Abstractions
- **Technology Trends**
- The Power Wall & Multi-processors
- Below Your Program
- Under the Covers
- Performance

The Computer Revolution

- Progress in computer technology
 - Underpinned by Moore's Law
- Makes novel applications feasible
 - Computers in automobiles
 - Cell phones
 - Human genome project
 - World Wide Web
 - Search Engines
- Computers are pervasive


Technology Scaling

In 1965, Intel's Gordon Moore predicted that the number of transistors that can be integrated on single chip would double about every two years

Number of transistors doubling per 18 months


What made this possible?

Number of transistors doubling per 24 months

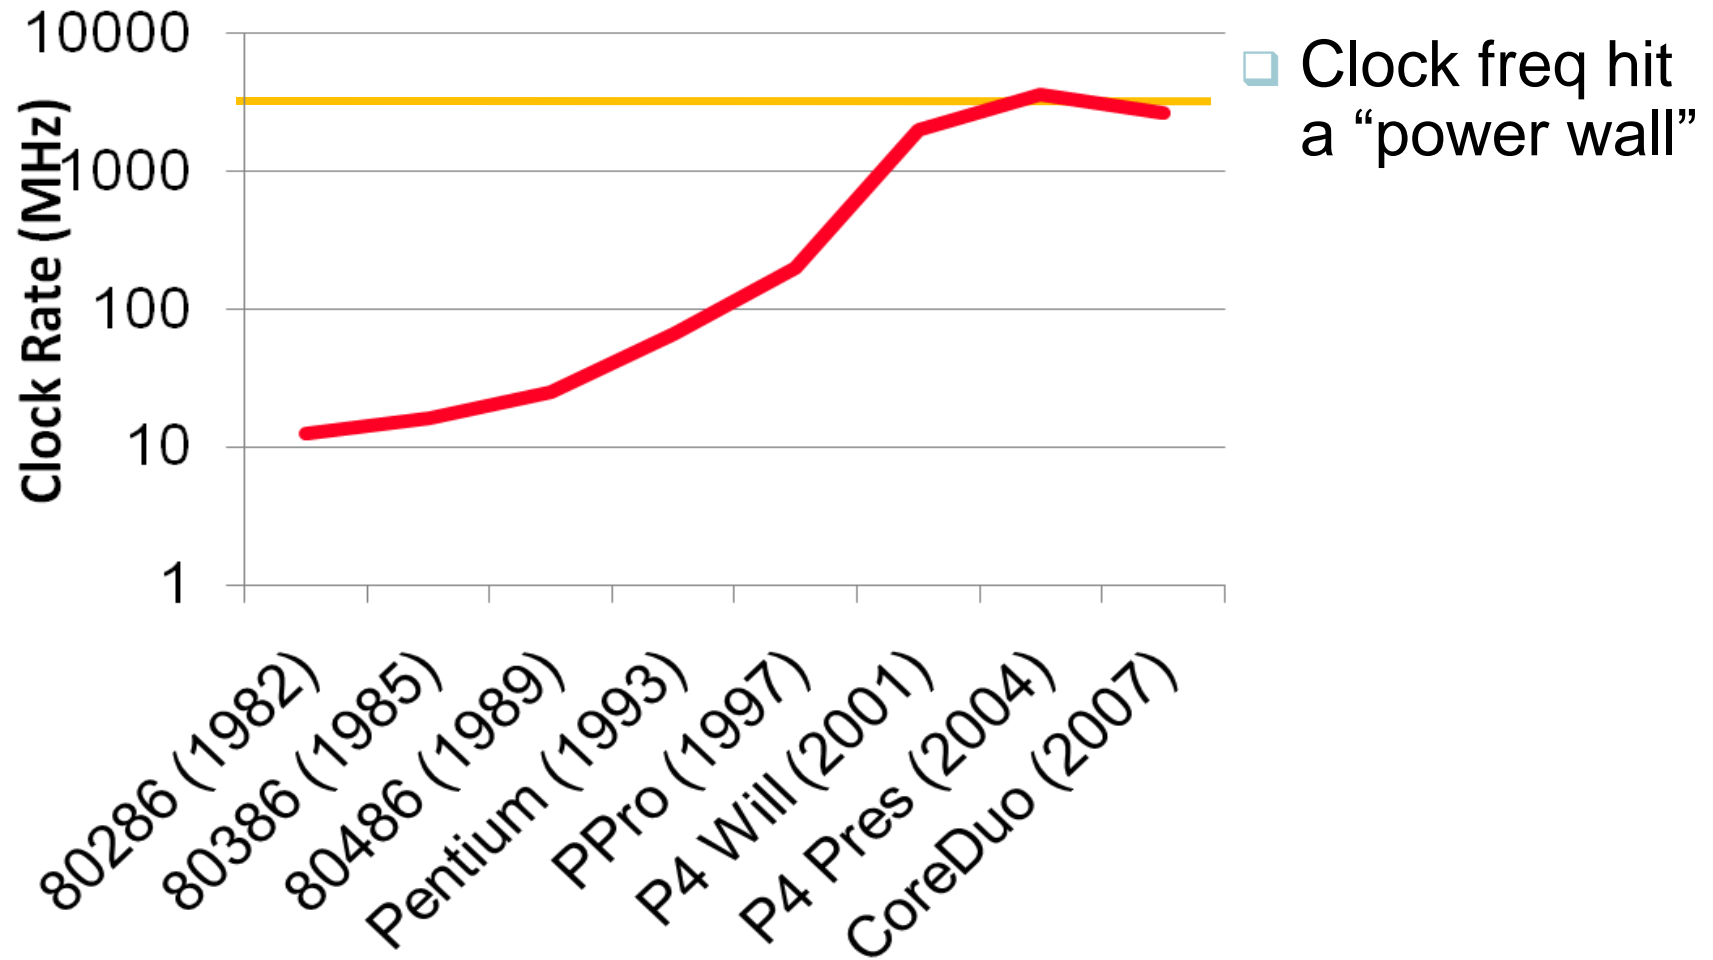


Technology Scaling Road Map (ITRS)

Year	2004	2006	2008	2010	2012
Feature size (nm)	90	65	45	32	22



Performance (Clock Freq) Scaling with Technology






CHAPTER 01

COMPUTER ABSTRACTIONS AND TECHNOLOGY

- Computer Abstractions
- Technology Trends
- **The Power Wall & Multi-processors**
- Below Your Program
- Under the Covers
- Performance

Power Trends (Power Wall)



- Heat dissipation? Over-clocking?

$$\text{Power}_{\text{Dynamic}} = \text{Capacitive load} \times \text{Voltage}^2 \times \text{Frequency}$$

↑
x30

↑
5V → 1V

↑
x300

Alternative ways to improve performance

- The power challenge has forced a change in the design of microprocessors
 - Since 2002 the rate of improvement in the response time of programs on desktop computers has slowed from a factor of 1.5 per year to less than a factor of 1.2 per year
- As of 2006 all desktop and server companies are shipping microprocessors with multiple processors – cores – per chip

Product	AMD Barcelona	Intel Nehalem	IBM Power 6	Sun Niagara 2
Cores per chip	4	4	2	8
Clock rate	2.5 GHz	~2.5 GHz	4.7 GHz	1.4 GHz
Power	120 W	~100 W	~100 W	94 W

- **Plan of record was to double the number of cores per chip per generation (about every two years). Did that happen?**

Multi-processors

- Multi-core microprocessors
 - More than one processor per chip
- Requires explicitly parallel programming
 - Compared to instruction level parallelism (ILP)
 - Hardware executes multiple instructions at once
 - Hidden from the programmer
 - Hard to do
 - Programming for performance
 - Load balancing
 - Optimizing communication and synchronization




CHAPTER 01

COMPUTER ABSTRACTIONS AND TECHNOLOGY

- Computer Abstractions
- Technology Trends
- The Power Wall & Multi-processors
- **Below Your Program**
- Under the Covers
- Performance

Below Your Program



- Application software
 - Written in high-level language
- Systems software
 - Compiler / Assembler: translate HLL code to machine code
 - Operating System: service code
 - Handling input/output
 - Managing memory and storage
 - Scheduling tasks & sharing resources
- Hardware
 - CPU, memory, I/O controllers

Levels of Program Code

- High-level language
 - Level of abstraction closer to problem domain
 - Provides for productivity and portability
- Assembly language
 - Textual representation of instructions
- Hardware representation
 - Binary digits (bits)
 - Encoded instructions and data

High-level
language
program
(in C)

```
swap(int v[], int k)
{int temp;
  temp = v[k];
  v[k] = v[k+1];
  v[k+1] = temp;
}
```

Compiler

Assembly
language
program
(for MIPS)

```
swap:
  muli $2, $5, 4
  add  $2, $4, $2
  lw   $15, 0($2)
  lw   $16, 4($2)
  sw   $16, 0($2)
  sw   $15, 4($2)
  jr   $31
```

Assembler

Binary machine
language
program
(for MIPS)

```
000000001010000100000000000011000
000000000000110000001100000100001
100011000110001000000000000000000
100011001111001000000000000000100
101011001111001000000000000000000
101011000110001000000000000000100
00000011111000000000000000001000
```




CHAPTER 01

COMPUTER ABSTRACTIONS AND TECHNOLOGY

- Computer Abstractions
- Technology Trends
- The Power Wall & Multi-processors
- Below Your Program
- Under the Covers
- Performance

What happens to the machine code...





- Inputs/outputs include
 - User-interface devices
 - Display, keyboard, mouse
 - Storage devices
 - Hard disk, CD/DVD, flash
 - Network adapters
 - For communicating with other computers

Inside the Processor (CPU)

- Datapath: performs operations on data
- Control: sequences datapath, memory, ...
- Cache memory
 - Small fast SRAM memory for immediate access to data

Inside the Processor

- AMD Barcelona: 4 processor cores





CHAPTER 01

COMPUTER ABSTRACTIONS AND TECHNOLOGY

- Computer Abstractions
- Technology Trends
- The Power Wall & Multi-processors
- Below Your Program
- Under the Covers
- Performance



UNDERSTANDING PERFORMANCE

Understanding Performance





- Algorithm
 - Determines number of operations executed
- Programming language, compiler, architecture
 - Determine number of machine instructions executed per operation
- Processor and memory system
 - Determine how fast instructions are executed
- I/O system (including OS)
 - Determines how fast I/O operations are executed

Why Know Performance?

- Measure, Report, and Summarize
- Make intelligent choices
- See through the marketing hype
- Key in understanding the underlying organizational motivation
 - Why is some hardware better than others for different programs?
 - What factors of system performance are hardware related?
(e.g., Do we need a new machine, or a new operating system?)
 - How does the machine's instruction set affect performance?

Defining Performance

- Which airplane has the best performance?



Response Time and Throughput

- Response time
 - How long it takes to do a task
- Throughput
 - Total work done per unit time
 - e.g., tasks/transactions/... per hour
- How are response time and throughput affected by
 - Replacing the processor with a faster version?
 - Adding more processors?
- We'll focus on response time for now...

Relative Performance

- Define Performance = $1/\text{Execution Time}$
- “X is n time faster than Y”

$$\begin{aligned} & \text{Performance}_X / \text{Performance}_Y \\ &= \text{Execution time}_Y / \text{Execution time}_X = n \end{aligned}$$


- Example: time taken to run a program
 - 10s on A, 15s on B
 - $\text{Execution Time}_B / \text{Execution Time}_A$
 $= 15\text{s} / 10\text{s} = 1.5$
 - So A is 1.5 times faster than B

Measuring Execution Time

- Elapsed time
 - Total response time, including all aspects
 - Processing, I/O, OS overhead, idle time
 - Determines system performance
- CPU time
 - Time spent processing a given job
 - Discounts I/O time, other jobs' shares
 - Comprises user CPU time and system CPU time
 - Different programs are affected differently by CPU and system performance

CPU Clocking

- Operation of digital hardware governed by a constant-rate clock



- Clock period: duration of a clock cycle
 - e.g., $250\text{ps} = 0.25\text{ns} = 250 \times 10^{-12}\text{s}$
- Clock frequency (rate): cycles per second
 - e.g., $4.0\text{GHz} = 4000\text{MHz} = 4.0 \times 10^9\text{Hz}$

CPU Time

$$\begin{aligned}\text{CPU Time} &= \text{CPU Clock Cycles} \times \text{Clock Cycle Time} \\ &= \frac{\text{CPU Clock Cycles}}{\text{Clock Rate}}\end{aligned}$$

- Performance improved by
 - Reducing number of clock cycles
 - Increasing clock rate
 - Hardware designer must often trade off clock rate against cycle count

CPU Time Example

- A program runs on computer A with a 2 GHz clock in 10 seconds. What clock rate must a computer B run at to run this program in 6 seconds? Unfortunately, to accomplish this, computer B will require 1.2 times as many clock cycles as computer A to run the program.

CPU Time Example

- Computer A: 2GHz clock, 10s CPU time
- Designing Computer B
 - Aim for 6s CPU time
 - Can do faster clock, but causes $1.2 \times$ clock cycles
- How fast must Computer B clock be?

$$\text{Clock Rate}_B = \frac{\text{Clock Cycles}_B}{\text{CPU Time}_B} = \frac{1.2 \times \text{Clock Cycles}_A}{6s}$$

$$\begin{aligned}\text{Clock Cycles}_A &= \text{CPU Time}_A \times \text{Clock Rate}_A \\ &= 10s \times 2\text{GHz} = 20 \times 10^9\end{aligned}$$

$$\text{Clock Rate}_B = \frac{1.2 \times 20 \times 10^9}{6s} = \frac{24 \times 10^9}{6s} = 4\text{GHz}$$

Exercise

- The same set of instructions are being executed on two CPUs A and B. A is faster than B. The reason could be
 1. A is having a complex circuit
 2. The clock frequency of B is less than A
 3. The clock frequency of B is higher than A
 4. The clock rate of A is less than B

Instruction Count and CPI

Clock Cycles = Instruction Count \times Cycles per Instruction

CPU Time = Instruction Count \times CPI \times Clock Cycle Time

$$= \frac{\text{Instruction Count} \times \text{CPI}}{\text{Clock Rate}}$$

- Instruction Count for a program
 - Determined by program, ISA and compiler
- Average cycles per instruction
 - Determined by CPU hardware
 - If different instructions have different CPI
 - Average CPI affected by instruction mix

Using the Performance Equation

- Computers A and B implement the same ISA. Computer A has a clock cycle time of 250ps and an effective CPI of 2.0 for some program and computer B has a clock cycle time of 500ps and an effective CPI of 1.2 for the same program. Which computer is faster and by how much?

CPI Example

- Computer A: Cycle Time = 250ps, CPI = 2.0
- Computer B: Cycle Time = 500ps, CPI = 1.2
- Same ISA
- Which is faster, and by how much?

$$\begin{aligned}\text{CPU Time}_A &= \text{Instruction Count} \times \text{CPI}_A \times \text{Cycle Time}_A \\ &= 1 \times 2.0 \times 250\text{ps} = 1 \times 500\text{ps} \end{aligned}$$

A is faster...

$$\begin{aligned}\text{CPU Time}_B &= \text{Instruction Count} \times \text{CPI}_B \times \text{Cycle Time}_B \\ &= 1 \times 1.2 \times 500\text{ps} = 1 \times 600\text{ps}\end{aligned}$$

$$\frac{\text{CPU Time}_B}{\text{CPU Time}_A} = \frac{1 \times 600\text{ps}}{1 \times 500\text{ps}} = 1.2$$

...by this much

CPI in More Detail

- If different instruction classes take different numbers of cycles

$$\text{Clock Cycles} = \sum_{i=1}^n (\text{CPI}_i \times \text{Instruction Count}_i)$$

- Weighted average CPI

$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}} = \sum_{i=1}^n \left(\text{CPI}_i \times \frac{\text{Instruction Count}_i}{\text{Instruction Count}} \right)$$

Relative frequency

CPI Example

- Alternative compiled code sequences using instructions in classes A, B, C

Class	A	B	C
CPI for class	1	2	3
IC in sequence 1	2	1	2
IC in sequence 2	4	1	1

- Sequence 1: IC = 5

- Clock Cycles
 $= 2 \times 1 + 1 \times 2 + 2 \times 3$
 $= 10$
- Avg. CPI = $10/5 = 2.0$

- Sequence 2: IC = 6

- Clock Cycles
 $= 4 \times 1 + 1 \times 2 + 1 \times 3$
 $= 9$
- Avg. CPI = $9/6 = 1.5$

Performance Summary

The BIG Picture

$$\text{CPU Time} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}}$$

- Performance depends on
 - Algorithm: affects IC, possibly CPI
 - Programming language: affects IC, CPI
 - Compiler: affects IC, CPI
 - Instruction set architecture: affects IC, CPI, T_c

Exercise

- The CPU time of a program **cannot** be given by
 - clock cycles / clock rate
 - (instructions/program) x (clock cycles/instruction) x (seconds/clock cycle)
 - instruction count * CPI * clock cycle time
 - instruction count * Average CPI / clock cycle time

Exercise

- A program runs on *computer1* with a 4GHz clock in 5.0 seconds. Calculate the clock rate of *computer2* that can finish this program in 2.5 seconds. Unfortunately, to accomplish this, *computer2* will require 1.5 times as many clock cycles as *computer1* to run the program.
 - 8 GHz
 - 12 GHz
 - The correct clock rate is not provided in answer a and b
 - Not enough data are given to calculate the clock rate

Exercise

- P_1 , P_2 are two different hardware implementations of the same ISA. P_1 and P_2 have 4 GHz and 6 GHz clock rates respectively.
- If the peak performance is defined as the fastest rate that a computer can execute any instruction sequence, what are the peak performance of P_1 and P_2
- If the instructions of a certain program P compiles into equally among the classes of instructions given, calculate how much faster it would be to execute this program in P_2 than in P_1 .

Class	CPI on P_1	CPI on P_2
A	1	2
B	2	2
C	3	2
D	4	4
E	3	4

Exercise

- Which of the following is **incorrect** about the performance computation of a CPU
 - If you have two CPUs with you, the best way to do a performance comparison is by measuring and comparing the execution time of appropriate applications on the CPUs.
 - The instructions of a CPU are divided into classes based on their CPI.
 - The CPI is calculated as an average number as individual instructions take a different number of clock cycles depending on which part of the program they are used.
 - A computer architect uses formulas and simulations to check the performance of a CPU as she does not have access to the CPUs they are building until they build them.

Pitfall: Amdahl's Law

- Improving an aspect of a computer and expecting a proportional improvement in overall performance

$$T_{\text{improved}} = \frac{T_{\text{affected}}}{\text{improvement factor}} + T_{\text{unaffected}}$$

- Example: multiply accounts for 80s/100s
 - How much improvement in multiply performance to get 5× overall?

$$20 = \frac{80}{n} + 20 \quad \quad \quad \blacksquare \text{ Can't be done!}$$

- Corollary: make the common case fast

Exercise

- According to Amdahl's law, the improvement in overall performance of a system can be given by
 - $T(\text{improved}) = (T(\text{unaffected}) / \text{improvement factor}) + T(\text{affected})$
 - $T(\text{improved}) = (T(\text{affected}) / \text{improvement factor}) + T(\text{unaffected})$
 - $T(\text{improved}) = (T(\text{affected}) * \text{improvement factor}) + T(\text{unaffected})$
 - $T(\text{improved}) = (T(\text{affected}) / \text{improvement factor}) - T(\text{unaffected})$

Concluding Remarks

- Cost/performance is improving
 - Due to underlying technology development
- Hierarchical layers of abstraction
 - In both hardware and software
- Instruction set architecture
 - The hardware/software interface
- Execution time: the best performance measure
- Power is a limiting factor
 - Use parallelism to improve performance